



Tokens per proposition (Sanskrit / English), 95% bootstrap CI

Deployed practice: every arm vs English o200k (200k, general domain)

Matched control: Sanskrit T1/T2 vs English E1 (same algorithm, vocab, training corpus); filled = pair-matched, hollow = byte-matched E1_bm; dotted = bytes



* provisional: trained on parallel-corpus training splits Not shown (unavailable this run): T3_indicsuper omitted (no candidate tokenizer could be loaded).