



Tokens per proposition by sentence length, matched control, 95% bootstrap CI

Binned on the ENGLISH side's word count — selects wordy denominators into the high bins



Binned on the SANSKRIT side's word count — mirror bias: wordy numerators into the high bins



— T1_bpe_raw_32k* / E1_bpe_32k
 — T2_unigram_raw_32k* / E1_unigram_32k
 - - T1_bpe_raw_32k* / E1_bpe_32k_bm
 - - T2_unigram_raw_32k* / E1_unigram_32k_bm

— T1_bpe_raw_64k* / E1_bpe_64k
 — T2_unigram_raw_64k* / E1_unigram_64k
 - - T1_bpe_raw_64k* / E1_bpe_64k_bm
 - - T2_unigram_raw_64k* / E1_unigram_64k_bm

top row bins by English whitespace word count, bottom row by Sanskrit whitespace word count; hollow = fewer than 30 pairs; each panel has its own y-axis, and a bin lying outside it is drawn as a triangle at that edge, annotated with the off-scale value