


Tokens per proposition vs the matched English control (E1), 95% bootstrap CI
raw-text arms (T1/T2) vs sandhi-split arms (T4), matched on algorithm and vocabulary



Provisional: every arm is trained on the Sanskrit (or English) side of two parallel corpora, not the monolingual corpus. split = sandhi + compound (samāsa) splitting, ByT5-Sanskrit, reconciled against the raw sentence. Splitter: chronbmm/sanskrit5-multitask@c0d2ada54f3d19903149425aa888a203601423f8.